



Department of Econometrics and Business Statistics

<http://monash.edu/business/ebs/research/publications>

Sparse Multiple Index Models for High-dimensional Nonparametric Forecasting

Nuwani Palihawadana, Rob J Hyndman, Xiaoqian Wang

February 2024

Working Paper no/yr



AACSB
ACCREDITED



Sparse Multiple Index Models for High-dimensional Nonparametric Forecasting

Nuwani Palihawadana

Department of Econometrics & Business Statistics
Clayton VIC 3800
Australia
Email: nuwani.kodikarapalihawadana@monash.edu
Corresponding author

Rob J Hyndman

Department of Econometrics & Business Statistics
Clayton VIC 3800
Australia
Email: rob.hyndman@monash.edu

Xiaoqian Wang

Department of Econometrics & Business Statistics
Clayton VIC 3800
Australia
Email: xiaoqian.wang@monash.edu

26 February 2024

JEL classification: C10,C14,C22

Sparse Multiple Index Models for High-dimensional Nonparametric Forecasting

Abstract

High-dimensionality is a common phenomenon in real-world forecasting problems. Oftentimes, forecasts are contingent on a long history of predictors, while the relationships between some predictors and the response of interest exhibit complex nonlinear patterns. In such a situation, a nonlinear “transfer function” model, with additivity constraints to mitigate the issue of *curse of dimensionality*, is a conspicuous choice. Particularly, nonparametric *additive index models* greatly reduce the number of parameters to be estimated in comparison to a general additive model. In this paper, we present a novel algorithm for estimating high-dimensional nonparametric additive index models, with simultaneous variable selection, which we call **SMI** (Sparse Multiple Index) **Model**. The SMI Model algorithm is based on an iterative procedure that applies mixed integer programming to solve an ℓ_0 -regularised nonlinear least squares problem. We demonstrate the functionality and the characteristics of the proposed algorithm through a simple simulation exercise. We also illustrate the use of the SMI Model algorithm in two empirical applications related to forecasting heat exposure related daily mortality and daily solar intensity.

Keywords: Additive index models, Variable selection, Dimension reduction, Mixed integer programming

1 Introduction

Forecasts are often contingent on a very long history of predictors. For example, when forecasting half-hourly electricity demand, it is common to use at least a week of historical half-hourly temperatures and other weather observations (Hyndman & Fan 2010). Similarly, when forecasting bore levels, rainfall data from up to thousand days earlier can impact the result (Bakker & Schaars 2019) due to the complex flow dynamics of rainfall into aquifers.

On the other hand, in most of these applications, the relationships between the predictors and the response variable exhibit complex nonlinear patterns. For instance, the relationship between

electricity demand and temperature is often nonlinear (Hyndman & Fan 2010; Fan & Hyndman 2012).

These examples suggest a possible nonlinear “*transfer function*” model of the form

$$y_t = f(\mathbf{x}_t, \mathbf{x}_{t-1}, \dots, \mathbf{x}_{t-p}, y_{t-1}, \dots, y_{t-k}) + \varepsilon_t, \quad (1)$$

where y_t is the observation of the response variable at time t , \mathbf{x}_t is a vector of predictors at time t , and ε_t is the random error. By including lagged values of y_t along with the lagged predictors, we allow for any serial correlation in the data. However, it makes the resulting function difficult to interpret. An alternative formulation is

$$y_t = f(\mathbf{x}_t, \mathbf{x}_{t-1}, \dots, \mathbf{x}_{t-p}) + g(\varepsilon_t, \varepsilon_{t-1}, \dots, \varepsilon_{t-k}),$$

which is more difficult to estimate, but makes it simpler to interpret the effect of the predictors on the response variable.

When applying the transfer function model to forecast lengthy time series with complex patterns, the form of f is nonlinear, involving complicated interactions, and with a high value of p .

Typically, the form of f involves many ad hoc model choices. It is essentially impossible to estimate a p -dimensional function for large p due to the curse of dimensionality (Bellman 1957; Stone 1982). Instead, we normally impose some form of additivity, along with some low-order interactions.

For example, Fan & Hyndman (2012) proposed a *semi-parametric additive model* to obtain short-term forecasts of the half-hourly electricity demand for power systems in the Australian National Electricity Market. In this model, f is assumed to be fully additive, and is used to capture the effects of recent predictor values on the demand. The main objective behind the use of this proposed semi-parametric model is to allow nonparametric components in a regression-based modelling framework with serially correlated errors (Fan & Hyndman 2012). The model fitted for each half-hourly period (q) can be written as

$$\log(y_{t,q}) = h_q(t) + f_q(w_{1,t}, w_{2,t}) + a_q(y_{t-p}) + \varepsilon_t,$$

where the response variable is the logarithm of electricity demand at time t (measured in the half-hourly intervals) during period q . The term $h_q(t)$ models several calendar effects that are included as linear terms. The temperature effects are modelled using the nonparametric

component $f_q(w_{1,t}, w_{2,t})$, while the nonparametric term $a_q(y_{t-p})$ captures the lagged effects of the response. It is important to notice here that the error term ε_t is serially uncorrelated in each half-hourly model, because the serial correlation is eliminated by the inclusion of the lagged responses in the model. However, there will still be some correlation between the residuals from the various half-hourly models (Fan & Hyndman 2012).

Similarly, a *distributed lag model* was proposed by Wood (2017) to forecast daily death rate in Chicago using measurements of several air pollutants. In this model, the response variable is modelled via a sum of smooth functions of lagged predictor variables, which is quite similar in nature to the semi-parametric additive model used by Fan & Hyndman (2012). However, unlike in Fan & Hyndman (2012), Wood (2017) suggested to allow the smooth functions for lags of the same covariate to vary smoothly over lags, preventing large differences in estimated effects between adjacent lags. Thus, the model is of the form

$$\log(y_t) = f_1(t) + \sum_{k=0}^K f_2(p_{t-k}, k) + \sum_{k=0}^K f_3(o_{t-k}, w_{t-k}, k),$$

where y_t is the death rate at day t , and f_1 is a nonparametric term to capture the *time* effect. The model incorporates the current value ($k = 0$) and several lagged values ($k = 1, \dots, K$) of the predictors, where the *distributed lag effect* of a single predictor variable, and of an interaction of two predictor variables are captured by the sum of nonparametric terms $\sum_{k=0}^K f_2(p_{t-k}, k)$ and $\sum_{k=0}^K f_3(o_{t-k}, w_{t-k}, k)$ respectively. The smooth functions f_2 and f_3 are proposed to be estimated using *tensor product smooths*.

For more examples, Ho, Chen & Hwang (2020) used semi-parametric additive models to estimate ground-level $PM_{2.5}$ concentrations in Taiwan, while nonparametric additive models were utilised by Ibrahim et al. (2022) for predicting census survey response rates. Furthermore, Ravindra et al. (2019) provided a comprehensive review of the applications of additive models for environmental data, with a special focus on air pollution, climate change, and human health related studies.

While such models have been used to address problems including electricity demand, air quality related mortality rate and groundwater level forecasting etc. (Fan & Hyndman 2012; Hyndman & Fan 2010; Wood 2017; Peterson & Western 2014; Rajaei, Ebrahimi & Nourani 2019), there are still a number of unresolved issues in their applications. In this paper, we attempt to address two of those issues. Firstly, even though nonparametric additive models act as a remedy to the curse of dimensionality as we discussed earlier, the estimation of the model is still challenging in a high-dimensional setting due to the large number of nonparametric components to be

estimated. Secondly, there is a noticeable subjectivity in the selection of predictor variables (from the available predictors) for the model, where in most of the applications of interest we discussed above, the predictor choices in the final model are mainly based on empirical explorations or domain expertise.

There are a number of previous studies that have attempted to address the issue of variable selection in nonparametric/semi-parametric additive models to some extent, using various techniques. For example, Huang, Horowitz & Wei (2010) used a *Least Absolute Shrinkage and Selection Operator (LASSO)* (Tibshirani 1996) based procedure for variable selection in nonparametric additive models, whereas Fan & Hyndman (2012) used a straightforward backward elimination technique to achieve selection. Moreover, Ibrahim et al. (2022) and Hazimeh, Mazumder & Radchenko (2023) used Mixed Integer Programming based methodologies to provide a solution to the *best subset selection* problem in nonparametric additive models. More details of these methods are discussed later in Section 2.

In this paper, however, we are interested in high-dimensional applications that exhibit complicated interactions among predictors (specially in the presence of large number of lagged variables), as well as correlated errors. In such a situation, “*index models*” (refer Section 2.2) seem to be useful for improving the flexibility of the broader class of nonparametric additive models (Radchenko 2015), while mitigating the difficulty of estimating a nonparametric component for each individual predictor.

To our knowledge, no previous research has been done to look at how the predictor choices can be made more objective and principled in nonparametric additive index models. Hence, our goal was to develop a methodology for optimal predictor selection in the context of high-dimensional nonparametric additive index models. Moreover, due to computational advancements in the field, the use of *Mathematical Optimisation* concepts in solving statistical problems has gained a lot of interest in the recent past (Theußl, Schwendinger & Hornik 2020). This motivated us to develop a variable selection algorithm based on mathematical optimisation techniques.

Additionally, it is crucial to point out that any such variable selection methodology naturally renders inferential statistics invalid, since we do not assume the resulting model obtained through the variable selection procedure to represent the true data generating process. Hence, our focus in this paper is only on improving forecasts, but not on making inferences on the resulting parameter estimates.

The rest of this paper is organised as follows. In Section 2, we provide a concise exposition of related ideas and previous work, while establishing the foundation for this paper. Section 3 presents our proposed model, *Sparse Multiple Index Model* (SMI Model), and describes the variable selection algorithm and estimation procedure. In Section 4, we demonstrate the functionality and the characteristics of the proposed algorithm through a simulation experiment. Section 5 illustrates two empirical applications of the proposed estimation and variable selection methodology, related to forecasting heat exposure related daily mortality and daily solar intensity. Concluding remarks are given in Section 6.

2 Background

2.1 Variable Selection in Nonparametric Additive Models

As discussed in Section 1, the estimation of nonparametric function f (Equation 1) becomes infeasible in high-dimensional settings (i.e. number of predictors is very large) due to curse of dimensionality. As a result, *nonparametric additive models* have been employed with growing popularity. Let $(y_i, x_i), i = 1, \dots, n$, be independent and identically distributed (i.i.d) observations, and $x_i = (x_{i1}, \dots, x_{ip})^T$ be a p -dimensional vector of predictor values. Then a nonparametric additive model can be written as

$$y_i = \sum_{j=1}^p f_j(x_{ij}) + \varepsilon_i, \quad i = 1, \dots, n, \quad (2)$$

where f_j 's are unknown functions (probably non-linear and smooth), and ε_i is the random error (Lian 2012). Even such an additivity condition is imposed, estimating the optimal predictive model will still be troublesome when p is very large (probably even larger than the sample size n) due to over-fitting (Lian 2012). Thus, it is natural to bring in the sparsity assumption, and assume that some of f_j 's are zero, which gives rise to the need of a variable selection method to differentiate between zero and non-zero components, while estimating the non-zero components (Huang, Horowitz & Wei 2010).

2.1.1 Backward Elimination

In the problem of forecasting long-term peak electricity demand, Hyndman & Fan (2010) used a stepwise procedure for variable selection through cross-validation. In the each half-hourly model fitted, the data is split into training and validation sets, and the predictors are selected into the model based on the Mean Squared Error (MSE) calculated for the validation set. Starting

from the full model, the predictive power of each variable is evaluated by dropping one at a time. A predictor, the removal of which contributed to a decrease in the validation MSE, is omitted from the model in subsequent steps (Hyndman & Fan 2010). Fan & Hyndman (2012) used a similar method except for the fact that they considered the Mean Absolute Percentage Error (MAPE) as the selection criterion. Therefore, both of these prior work use stepwise variable selection methodology based on out-of-sample forecasting performance.

2.1.2 Penalisation Methods

According to Huang, Horowitz & Wei (2010), there are numerous penalised methods for variable selection and parameter estimation in high-dimensional settings, including the *bridge estimator* proposed by Frank & Friedman (1993), the *Least Absolute Shrinkage and Selection Operator* (LASSO) by Tibshirani (1996), the *Smoothly Clipped Absolute Deviation Penalty* (SCAD) by Fan & Li (2001), and the *Minimum Concave Penalty* (MCP) by Zhang (2010). Among them, we observe that the LASSO and the SCAD penalties are appearing popularly in literature.

Tibshirani (1996) introduced the regularisation method, *LASSO*, for estimating linear models, which minimises the sum of squared residuals subject to the ℓ_1 penalty on the coefficients. Assume the classical linear regression model $y_i = \sum_{j=1}^p \beta_j x_{ij} + \varepsilon_i$, fitted for the data (y_i, \mathbf{x}_i) , $i = 1, \dots, n$, where y_i is the response, $\mathbf{x}_i = (x_{i1}, \dots, x_{ip})^T$ is a p -dimensional vector of predictors, $\boldsymbol{\beta} = (\beta_1, \dots, \beta_p)^T$ is the parameter vector corresponding to \mathbf{x}_i , and ε_i is the random error. Then, the LASSO estimator, $\hat{\boldsymbol{\beta}}_{LASSO}$, can be obtained by

$$\hat{\boldsymbol{\beta}}_{LASSO} = \min_{\boldsymbol{\beta}} \left\{ \left\| \mathbf{y} - \sum_{j=1}^p \mathbf{x}_j \beta_j \right\|_2^2 + \lambda \sum_{j=1}^p |\beta_j| \right\},$$

where $\mathbf{x}_j = (x_{1j}, \dots, x_{nj})^T$, and λ is a non-negative tuning parameter. The LASSO estimator reduces to the Ordinary Least Squares (OLS) estimator if λ is equal to zero (Konzen & Ziegelmann 2016). Due to the nature of the penalty applied, LASSO shrinks some of the coefficients towards zero, and sets the others exactly to zero, where the estimation of coefficients and variable selection are performed simultaneously (Konzen & Ziegelmann 2016).

While showing that the LASSO is not consistent for variable selection in certain situations, Zou (2006) introduced *Adaptive Lasso* (popularly known as “adaLASSO”); an extension of the LASSO method, which uses adaptive weights to penalise coefficients using the LASSO (i.e. ℓ_1)

penalty. Thus, the adaLASSO objective function can be written as

$$\hat{\beta}_{adaLASSO} = \min_{\beta} \left\{ \left\| \mathbf{y} - \sum_{j=1}^p \mathbf{x}_j \beta_j \right\|_2^2 + \lambda \sum_{j=1}^p w_j |\beta_j| \right\},$$

where the vector of weights $\mathbf{w} = (w_1, \dots, w_p)^T$ is estimated by $\hat{\mathbf{w}} = 1/|\hat{\beta}|^\gamma$ for $\gamma > 0$, which is a tuning parameter, and $\hat{\beta}$ being any consistent estimator of β (Zou 2006).

Yuan & Lin (2006) considered the problem of selecting groups of variables, and discussed extensions of three variable selection and estimation methods namely, *LASSO* (Tibshirani 1996), *Least Angle Regression Selection* (LARS, Efron et al. 2004), and *Non-negative Garrotte* (Breiman 1995). Consider an n -dimensional response vector \mathbf{y} , and an $n \times p$ matrix of predictor values \mathbf{X} . Then the **Group Lasso** estimator of the coefficients vector β is obtained by minimising

$$\frac{1}{2} \left\| \mathbf{y} - \sum_{\ell=1}^L \mathbf{X}_\ell \beta_\ell \right\|_2^2 + \lambda \sum_{\ell=1}^L \|\beta_\ell\|_{\mathbf{K}_\ell},$$

where \mathbf{X}_ℓ is an $n \times p_\ell$ sub-matrix in \mathbf{X} that corresponds to the ℓ^{th} group of predictors (p_ℓ is the number of predictors in ℓ^{th} group), β_ℓ is the corresponding vector of coefficients, $\ell = 1, \dots, L$, $\|\beta_\ell\|_{\mathbf{K}_\ell} = (\beta_\ell' \mathbf{K}_\ell \beta_\ell)^{\frac{1}{2}}$ with $\mathbf{K}_1, \dots, \mathbf{K}_L$ being a set of given positive definite matrices, and λ is a non-negative tuning parameter. Moreover, Simon et al. (2013) proposed **Sparse-Group Lasso**, which is a convex combination of general Lasso and Group Lasso methods, where the focus is on both “groupwise sparsity” (the number of groups with at least one nonzero coefficient), and “within group sparsity” (the number of nonzero coefficients within each nonzero group).

According to Fan & Li (2001), a penalty function used in penalised least squares approaches should have three properties. Firstly, it should be singular at origin to generate a solution that is sparse. Secondly, it should fulfill certain conditions to be stable in model selection. Finally, it should be able to generate unbiased estimates for large coefficients via being bounded by a constant. They argued that all those three conditions are not satisfied by the penalisation methods such as the bridge regression (Frank & Friedman 1993) and the LASSO (Tibshirani 1996). Hence they proposed the **SCAD** penalty function, which is defined in terms of its first derivative as

$$p'_\lambda(\theta) = \lambda \left\{ I(\theta \leq \lambda) + \frac{(a\lambda - \theta)_+}{(a-1)\lambda} I(\theta > \lambda) \right\},$$

for some $a > 2$, and $\theta > 0$ (Fan & Li 2001). According to Fan & Li (2001), the SCAD penalty function retains the favourable properties of both best subset selection and ridge regression, while having all three desired features, i.e., sparsity, stability, and unbiasedness.

Based on the above penalisation methods that are originally developed for linear models, Huang, Horowitz & Wei (2010) proposed a new penalisation method for variable selection in nonparametric additive model (Equation 2), named *Adaptive Group Lasso*. They approximated f_j 's using normalised B-spline bases, so that a linear combination of B-spline basis functions is used to represent an individual nonparametric component f_j . The proposed method is a generalisation of Adaptive Lasso method (Zou 2006) to the Group Lasso method (Yuan & Lin 2006).

When the nonparametric additive model in Equation 2 is considered, an obvious possibility is that some of the additive components (i.e. f_j 's) are being linear. For example, recall the electricity demand forecasting problem (Hyndman & Fan 2010; Fan & Hyndman 2012), where some of the calendar effects are included into the model as linear variables, whereas lagged temperature and lagged demand variables are included using nonlinear additive components. Such situations suggest the use of *semi-parametric partially linear additive models* that can be mathematically represented as

$$y_i = \sum_{j=1}^p f_j(x_{ij}) + \sum_{k=1}^q w_{ik}\beta_k + \varepsilon_i, \quad i = 1, \dots, n,$$

where x_j 's, $j = 1, \dots, p$, are a set of predictors that enter the model as nonparametric components, whereas w_k 's, $k = 1, \dots, q$ are another set of predictors that are included as linear components. While several studies have assumed that the number of nonparametric components are fixed, and performed variable selection only among the linear components of the model (Lian 2012; Guo et al. 2013; Liu, Wang & Liang 2011), Wang et al. (2014) introduced a methodology for selecting both linear and nonlinear components simultaneously, in the context of correlated, longitudinal data. They proposed the use of a *Penalised Quadratic Inference Function (PQIF) with double SCAD penalties* for variable selection and model estimation, where the correlation structure of the data was incorporated into the estimation method (see Wang et al. (2014) for details).

2.1.3 Time Series Aspect

It is worthwhile to briefly mention that there are extensions of the penalisation methods discussed above, which have specifically proposed to take the autocorrelation and lag structures in time series data into account.

Wang, Guodong & Tsai (2007) proposed an extension of the LASSO method for Regression with Autoregressive Error (REGAR) models. Park & Sakaori (2013) and Konzen & Ziegelmann (2016) proposed modifications to Adaptive Lasso method to incorporate the lag structures presented

in Autoregressive Distributed Lag (ADL) models into the variable selection and estimation methodology. The *Ordered Lasso* was introduced by Tibshirani & Suo (2016) to deal with time-lagged regression problems, where we forecast the response value at time t using the predictor values from K previous time points, assuming that the magnitude of regression coefficients decreases as the lagged predictor moves away from time t .

However, it is important to note that all the models considered in the above time series related work are linear; none of them include nonparametric terms.

2.2 Index Models

2.2.1 Single Index Model

The nonparametric additive model (Equation 2) estimates the relationship between the response and the predictors using a sum of univariate nonlinear functions corresponding to each individual predictor variable. Hence, it is incapable of handling the interactions among the predictors, which are ubiquitous in real-world problems (Zhang et al. 2008).

As a remedy, the *Single Index Model*, a generalisation of the linear regression model where the linear predictor is replaced by a semi-parametric component, is popularly being used in the literature (Radchenko 2015). Let y_i be the response, and \mathbf{x}_i be a p -dimensional predictor vector. Then the single index model can be written as

$$y_i = g\left(\boldsymbol{\alpha}^T \mathbf{x}_i\right) + \varepsilon_i, \quad i = 1, \dots, n,$$

where $\boldsymbol{\alpha}$ is a p -dimensional vector of unknown coefficients (i.e. parameters), g is an unknown univariate function, and ε_i is the random error (Stoker 1986; Härdle, Hall & Ichimura 1993). The linear combination $\boldsymbol{\alpha}^T \mathbf{x}_i$ is called the *index*. Single index model is viewed as a viable alternative to the additive model since it offers more flexibility and interpretability (Radchenko 2015).

According to Radchenko (2015), single index models have widely been used in scenarios with fairly low and moderate dimensionality, where the corresponding estimation and variable selection techniques are not directly applicable to the high-dimensional setting. The error sum of squares of the model being non-convex with respect to index coefficients, is the main reason behind the existence of very limited number of methods in high-dimensional case (Radchenko 2015). For an extensive summary of available methods, we refer to Radchenko (2015).

2.2.2 Multiple Index Models

Projection Pursuit Regression

Friedman & Stuetzle (1981) introduced *Projection Pursuit Regression (PPR)* by extending the nonparametric additive model (Equation 2) to enable the modelling of interactions among predictor variables. On the other hand, PPR is an extension of the single index model to an “additive index model”, given by

$$y_i = \sum_{j=1}^q g_j(\alpha_j^T x_i) + \varepsilon_i, \quad i = 1, \dots, n,$$

where y_i is the response, x_i is a p -dimensional predictor vector, $\alpha_j = (\alpha_{j1}, \dots, \alpha_{jp})^T, j = 1, \dots, q$ are p -dimensional projection vectors (or vectors of “index coefficients”), g_j ’s are unknown univariate functions, and ε_i is the random error.

Instead of estimating a single index, PPR estimates multiple indices and connects them to the response through a sum of univariate nonlinear functions. These indices are constructed through a *Projection Pursuit (PP)* (Kruskal 1969; Friedman & Tukey 1974) algorithm, which is considered to be “interesting” low-dimensional projections of a high-dimensional feature space, obtained through the maximisation of an appropriate objective function or a “projection index” (Huber 1985).

According to Zhang et al. (2008), PPR increases the power of additive models in high-dimensional settings, but it has two major drawbacks. Firstly, since PP increases the freedom of the additive model, it tends to overfit in a situation, where there are a lot of unimportant predictors. Secondly, the interpretation of the model estimated by PPR will be troublesome as many non-zero elements will be present in each projection vector α_j . To overcome these issues, Zhang et al. (2008) introduced an ℓ_1 regularised projection pursuit algorithm, where the resultant regression model is named as *Sparse Projection Pursuit Regression (SpPPR)*. In SpPPR, an ℓ_1 penalty (i.e. a LASSO penalty) on index coefficients is added to the cost function (the squared error) at each iteration of the PP, thereby performing variable selection and model estimation simultaneously. See Zhang et al. (2008) for more details.

Although Zhang et al. (2008) claimed that the SpPPR algorithm can detect important predictors even in a noisy data set, our experiments show that it is not particularly scalable for large data sets with both higher number of predictors and observations.

Group-wise Additive Index Model

Even though PPR introduces flexibility and the ability to model interactions among predictors into additive models, the indices obtained through PPR contain all the predictors at hand. Hence, even with a variable selection mechanism like SpPPR (Zhang et al. 2008), PPR creates indices possibly by mixing heterogeneous variables in a single linear combination, making very little sense in terms of interpretability (Masselet et al. 2022).

Typically, in many real-world problems, natural groupings can be identified in predictor variables. For example, naturally interacting variables can be grouped together, such as several lags of a predictor, weather related variables, and genes or proteins that are grouped by biological pathways in a biological study (Masselet et al. 2022; Wang, Xu & Zhu 2015).

This suggests the use of a *Group-wise Additive Index Model (GAIM)*, which can be written as

$$y_i = \sum_{j=1}^p g_j(\alpha_j^T x_{ij}) + \varepsilon_i, \quad i = 1, \dots, n,$$

where y_i is the univariate response, $x_{ij} \in \mathbb{R}^{l_j}$, $j = 1, \dots, p$ are naturally occurring p groups of predictors, which are p non-overlapping subsets of x_i - the vector of all predictors, α_j is a l_j -dimensional vector of index coefficients corresponding to the index $h_{ij} = \alpha_j^T x_{ij}$, g_j is an unknown (possibly nonlinear) component function, and ε_i is the random error, which is independent of x_i (Wang et al. 2015; Masselet et al. 2022).

Since GAIM uses groups of predictors that are naturally or logically belonging together to construct indices, such derived indices will be more expressive and interpretable. However, at the same time, this introduces a certain level of subjectivity into the model formulation as different users can group the available predictors in different ways based on different logical reasoning.

In this paper, our aim is to reduce that subjectivity induced by personal judgment or domain expertise. Hence, we propose a methodology that injects more objectivity into the estimation of multiple index models by algorithmically grouping predictors into indices, resulting in a model with a higher predictive accuracy.

Constrained Group-wise Additive Index Model

The *Constrained Group-wise Additive Index Model (CGAIM)* was proposed by Masselet et al. (2022) for constructing comprehensive and easily interpretable indices from a large set of

explanatory variables. The model of interest is a *semi-parametric group-wise additive index model* given by

$$y_i = \beta_0 + \sum_{j=1}^p g_j(\alpha_j^T x_{ij}) + \sum_{k=1}^d f_k(w_{ik}) + \theta^T u_i + \varepsilon_i, \quad i = 1, \dots, n,$$

where y_i is the univariate response, β_0 is the model intercept, $x_{ij} \in \mathbb{R}^{l_j}$, $j = 1, \dots, p$ are naturally occurring p groups of predictor vectors (i.e. it is assumed that the predictor groupings are known in advance), which are p subsets of x_i - the q -dimensional vector of all predictors entering indices, α_j is the vector of index coefficients corresponding to the index $h_{ij} = \alpha_j^T x_{ij}$, and g_j is the corresponding nonlinear link function (possibly estimated by a spline). The additional predictor variables that are helpful in predicting y_i , but do not enter any of the indices are two-fold: a covariate that relates to the response through a nonlinear function f_k , denoted by w_{ik} , and the vector of linear covariates denoted by u_i .

This is an extension of the GAIM that allows to impose constraints on the index coefficients as well as on the nonlinear link functions. In CGAIM, linear constraints of the form $C_j \alpha_j \geq 0$ can be imposed on the index coefficients, where $C_j \in \mathbb{R}^{d_j \times l_j}$, and d_j is the number of constraints. Moreover, shape constraints such as monotonicity, convexity or concavity can be imposed on the nonparametric functions. This modification allows to incorporate prior knowledge or operational requirements into the model estimation.

First, considering only the additive index part of the model, and given $(y_i, x_{i1}, \dots, x_{iq})$, $i = 1, \dots, n$ be the observed data, where the q predictors are grouped into p groups, the estimation problem of the CGAIM can be formulated as

$$\begin{aligned} \min_{\alpha, \beta_0} \quad & \sum_{i=1}^n \left[y_i - \beta_0 - \sum_{j=1}^p g_j(\alpha_j^T x_{ij}) \right]^2, \\ \text{s.t.} \quad & C\alpha \geq 0, \quad g_j \in m, \end{aligned} \tag{3}$$

where $\alpha = [\alpha_1^T, \dots, \alpha_p^T]^T$, β_0 is the model intercept, $C \in \mathbb{R}^{d \times q}$, d is the number of constraints on the index coefficients vector α , and m is a shape constraint imposed on g_j (Massetot et al. 2022).

Notice that α_j s behave non-linearly in Equation 3, and hence, this is a non-linear least squares problem. Accordingly, Masselot et al. (2022) introduced an efficient iterative algorithm for estimating the CGAIM based on *Sequential Quadratic Programming* (SQP), one of the most successful techniques for solving nonlinear constrained optimisation problems (Boggs & Tolle 1995). For details of the CGAIM algorithm refer to Masselot et al. (2022).

2.3 Mathematical Optimisation for Variable Selection

2.3.1 Mathematical Optimisation

Optimisation plays a major role in both decision science and physical systems evaluation. *Mathematical Optimisation* or *Mathematical Programming* can be defined as the minimisation (or maximisation) of a function subject to restrictions on the unknowns/parameters of that function (Nocedal & Wright 2006). Hence, a mathematical optimisation problem can be written as

$$\begin{aligned} \min_{\mathbf{x}} \quad & f_0(\mathbf{x}) \\ \text{s.t.} \quad & f_i(\mathbf{x}) \leq b_i, \quad i = 1, \dots, m \end{aligned} \tag{4}$$

where the vector of unknowns or parameters of the problem is given by $\mathbf{x} = (x_1, \dots, x_n)^T$, the *objective function* is denoted by $f_0 : \mathbb{R}^n \rightarrow \mathbb{R}$, the *constraint functions* are given by $f_i : \mathbb{R}^n \rightarrow \mathbb{R}$, $i = 1, \dots, m$, and the bounds of the constraints are denoted by $\mathbf{b} = (b_1, \dots, b_m)^T$. A vector of values \mathbf{x}^* that results in the smallest value for the objective function among all vectors that satisfy the stated constraints, is called the *optimal* value or the *solution* to the problem (Boyd & Vandenberghe 2004). After mathematically formulating the optimisation problem as above (Equation 4), an appropriate *optimisation algorithm* is used to obtain the solution \mathbf{x}^* (Nocedal & Wright 2006).

Based on the form of the objective function and the constraints, various types of optimisation problems are identified.

An optimisation problem is known as a **Linear Program** (LP) when both the objective function and the constraints in Equation 4 (i.e. all $f_i, i = 0, \dots, m$) are linear. Hence, a LP can be written as

$$\begin{aligned} \min_{\mathbf{x}} \quad & \mathbf{a}_0^T \mathbf{x} \\ \text{s.t.} \quad & \mathbf{A}\mathbf{x} \leq \mathbf{b}, \end{aligned} \tag{5}$$

where \mathbf{x} is the vector that contains the parameters to be optimised, and $\mathbf{a}_0 \in \mathbb{R}^n$ is the vector of coefficients of the objective function. The matrix of coefficients in the constraints is denoted by $\mathbf{A} \in \mathbb{R}^{m \times n}$, and \mathbf{b} is the vector containing the upper bounds of the constraints. All LPs are *convex* optimisation problems (Theußl, Schwendinger & Hornik 2020).

The LP problem given in Equation 5 can be generalised to involve a quadratic term in the objective function, in which case it is called a **Quadratic Program** (QP). A QP can be written as

$$\begin{aligned} \min_x \quad & \frac{1}{2}x^T Q_0 x + a_0^T x \\ \text{s.t.} \quad & Ax \leq b, \end{aligned}$$

where $Q_0 \in \mathbb{R}^{n \times n}$. Unless the matrix Q_0 is positive semi-definite, a QP is non-convex (Theußl, Schwendinger & Hornik 2020).

If a linear objective function is minimised over a *convex cone*, such an optimisation problem is called a **Conic Program** (CP), which can be written as

$$\begin{aligned} \min_x \quad & a_0^T x \\ \text{s.t.} \quad & Ax + s = b, \quad s \in \mathcal{K}, \end{aligned}$$

where \mathcal{K} denotes a nonempty closed convex cone. CPs are designed to model convex optimisation problems (Theußl, Schwendinger & Hornik 2020).

If we restrict some of the unknowns/parameters in an optimisation problem to take only integer values, then that optimisation problem is called a **Mixed Integer Program** (MIP). For example, if we constraint $x_k \in \mathbb{Z}$ for at least one $k, k \in \{1, \dots, n\}$ in the optimisation problem given by Equation 4, then the optimisation problem becomes a MIP. If all the unknowns of an optimisation problem are constrained to be integers, such a problem is referred to as a pure **Integer Program** (IP), whereas if all the unknowns are bounded between zero and one (i.e. $x \in \{0, 1\}^n$), the optimisation problem is referred to as a **Binary (Integer) Program** (Theußl, Schwendinger & Hornik 2020). MIPs are hard to solve as they are non-convex due to the integer constraints. However, a growth in the number of commercial as well as non-commercial MIP solvers has made it possible to solve MIP problems conveniently and directly.

2.3.2 Variable Selection

Mathematical optimisation is fundamentally important in statistics, as many statistical problems including regression, classification, and other types of estimation/approximation problems can be re-interpreted as optimisation problems (Theußl, Schwendinger & Hornik 2020). Thus, the problem of variable selection - one of the prolonged interests of statisticians, has also benefited from using optimisation concepts, particularly MIP and convex optimisation, in the recent past.

For example, Bertsimas, King & Mazumder (2016) used a mixed integer optimisation procedure to solve the classical best subset selection problem in a linear regression. They developed a discrete optimisation method by extending modern first-order continuous optimisation techniques. The method can produce near-optimal solutions that would serve as warm starts for a MIP algorithm, which would choose the best k features out of p predictors. Similarly, Hazimeh & Mazumder (2020) developed fast and efficient algorithms based on coordinate descent and local combinatorial optimisation to solve the same best subset selection (or ℓ_0 -regularised least squares) problem through re-formulating local combinatorial search problems as structured MIPs.

Furthermore, Hazimeh, Mazumder & Radchenko (2023) proposed a group-wise variable selection methodology, based on discrete mathematical optimisation, which is applicable to both ℓ_0 -regularised linear regression and nonparametric additive models in a high-dimensional setting. They formulated the group ℓ_0 -based estimation problem as a *Mixed Integer Second Order Cone Program (MISOCP)*, and proposed a new customised Branch-and-Bound (BnB) algorithm (Land & Doig 1960; Little et al. 1963) to obtain the global optimal solution to the MISOCP.

Through the study of above literature, we noticed that the mathematical optimisation based algorithms reduce computational cost of variable selection procedures in high-dimensional settings. This is largely due to the availability of efficient commercial solvers such as *Gurobi* and *CPLEX*. This motivated us to focus on a mathematical optimisation based procedure for developing our variable selection methodology.

3 Sparse Multiple Index Model

In this section, we develop a *Sparse Multiple Index Model* (hereafter referred to as SMI Model) to establish an objective and a principled methodology for estimating high-dimensional nonparametric additive index models, with optimal predictor selection.

3.1 Model

The model of interest is a semi-parametric additive index model, which can be written as

$$y_i = \beta_0 + \sum_{j=1}^p g_j(\alpha_j^T x_{ij}) + \sum_{k=1}^d f_k(w_{ik}) + \theta^T \mathbf{u}_i + \varepsilon_i, \quad i = 1, \dots, n, \quad (6)$$

where y_i is the univariate response, β_0 is the model intercept, $x_{ij} \in \mathbb{R}^{l_j}$, $j = 1, \dots, p$ are p subsets of x_i - the q -dimensional vector of all predictors entering indices, α_j is the l_j -dimensional vector of index coefficients corresponding to the index $h_{ij} = \alpha_j^T x_{ij}$, and g_j is the corresponding nonlinear link function (possibly estimated by a spline).

Based on the above model formulation, we make three main assumptions as follows to define the **SMI Model**:

1. The number of indices (i.e. the number of subsets of predictors) p is unknown, and will be estimated through the proposed algorithm;
2. The split of the predictors among indices is unknown, and will be determined by the proposed algorithm; and
3. A predictor variable (that is entering indices) can only enter one index (i.e. overlapping of predictors among indices is not allowed).

These assumptions further imply that the index coefficient vectors α_j s and the corresponding nonlinear link functions g_j s are also unknown, and will be estimated through the proposed algorithm.

Most importantly, we allow the possibility for the index coefficient vectors α_j s to have elements that are equal to zero, so that the predictors corresponding to such zero coefficients are dropped out from the model, achieving variable selection.

Moreover, it is important to note here that the possible number of indices in a SMI model ranges from 1 (i.e. all q predictors are in a single index) to q (i.e. each predictor is in a separate index). In other words, both the Single Index Model and the Additive Model are special cases of SMI Model.

In addition to the predictor variables that are entering indices, we also allow for predictors that do not enter any of the indices. These additional predictors are two-fold: a covariate that relates to the response through nonlinear function f_k denoted by w_{ik} , $k = 1, \dots, d$, and the vector of linear covariates denoted by u_i .

3.2 Optimisation Problem Formulation

Let q be the total number of predictors entering p non-overlapping subsets of size l_j , $j = 1, \dots, p$ (i.e. $\sum_{j=1}^p l_j = q$). The algorithm discussed in this paper apply to the SMI Model estimator given

below, where the squared error of the model (Equation 6) is minimised together with an ℓ_0 penalty term and an ℓ_2 (ridge) penalty term:

$$\begin{aligned} \min_{p, \alpha, g, \beta_0} \quad & \sum_{i=1}^n \left[y_i - \beta_0 - \sum_{j=1}^p g_j (\alpha_j^T x_{ij}) - \sum_{k=1}^d f_k(w_{ik}) - \theta^T u_i \right]^2 \\ & + \lambda_0 \sum_{j=1}^p \sum_{m=1}^{l_j} \mathbf{1}(\alpha_{jm} \neq 0) + \lambda_2 \sum_{j=1}^p \|\alpha_j\|_2^2 \end{aligned} \quad (7)$$

where $\alpha = [\alpha_1^T, \dots, \alpha_p^T]^T$, $g = \{g_1, g_2, \dots, g_p\}$, $\mathbf{1}(\cdot)$ is the indicator function, $\lambda_0 > 0$ is a tuning parameter that controls the number of selected predictors, and $\lambda_2 \geq 0$ is another tuning parameter that controls the strength of the additional shrinkage imposed on the estimated index coefficients.

Applying an ℓ_2 -penalty in addition to the ℓ_0 -penalty is motivated by related literature (Hazimeh & Mazumder 2020; Mazumder, Radchenko & Dedieuc 2022; Hazimeh, Mazumder & Radchenko 2023), where it is suggested that the prediction performance of best-subset selection is enhanced by the inclusion of an additional ridge penalty, especially when a low signal-to-noise ratio (SNR) is present.

3.3 MIQP Formulation

To solve the optimisation problem (Equation 7), we present a big-M based Mixed Integer Quadratic Program (MIQP) formulation:

$$\begin{aligned} \min_{p, \alpha, g, \beta_0, z} \quad & \sum_{i=1}^n \left[y_i - \beta_0 - \sum_{j=1}^p g_j (\alpha_j^T x_i) - \sum_{k=1}^d f_k(w_{ik}) - \theta^T u_i \right]^2 + \lambda_0 \sum_{j=1}^p \sum_{m=1}^q z_{jm} + \lambda_2 \sum_{j=1}^p \sum_{m=1}^q \alpha_{jm}^2 \\ \text{s.t.} \quad & |\alpha_{jm}| \leq M z_{jm} \quad \forall j, \forall m, \\ & \sum_{j=1}^p z_{jm} \leq 1 \quad \forall m, \\ & z_{jm} \in \{0, 1\}, \\ & j = 1, \dots, p, \quad m = 1, \dots, q, \end{aligned} \quad (8)$$

where p is the (unknown) number of indices, x_i is the q -dimensional vector of all predictors entering indices, $\alpha = [\alpha_1^T, \dots, \alpha_p^T]^T$, $g = \{g_1, g_2, \dots, g_p\}$, and $z = (z_1^T, \dots, z_p^T)^T$, $z_j = (z_{j1}, \dots, z_{jq})^T$, $j = 1, \dots, p$ such that $z_{jm} \in \{0, 1\}$, $m = 1, \dots, q$ for all j . In other words, we introduce a binary (i.e. indicator) variable corresponding to each predictor in each index. A pre-specified *big-M parameter* is denoted by $M < \infty$, and it should be sufficiently large. If α^* is an optimal solution

to the problem given in Equation 8, then the big-M parameter should satisfy $\max \left(|\alpha_{jm}^*| \right) \leq M$, where $j \in \{1, \dots, p\}$, and $m \in \{1, \dots, q\}$.

Notice that, here we formulate the MIQP to include all q predictors in each index so that in this case, α_j is a q -dimensional vector of index coefficients. However at the same time, as mentioned earlier, we introduce a set of binary variables corresponding to each predictor in each index, which serves two main purposes: firstly, these binary variables are used to make the “on-or-off” decisions of the predictors in the model; secondly, they contribute to decide which predictors belong to which index.

To further elaborate, first, the big-M constraints ensure that α_{jm} is zero if and only if z_{jm} is zero, and if $z_{jm} = 1$, then $|\alpha_{jm}| \leq M$. At the same time, the ℓ_0 -penalty term $\lambda_0 \sum_{j=1}^p \sum_{m=1}^q z_{jm}$ influences some of the binary variables z_{jm} to be zero, while the ℓ_2 -penalty term $\lambda_2 \sum_{j=1}^p \sum_{m=1}^q \alpha_{jm}^2$ enforces additional shrinkage on the estimated coefficients. Therefore, these components together perform a variable selection.

Next, when the set of binary variables $\mathbf{Z}_m = \{z_{1m}, z_{2m}, \dots, z_{pm}\}$ corresponding to the m^{th} , $m = 1, \dots, q$, predictor in all p indices is considered, according to the constraint $\sum_{j=1}^p z_{jm} \leq 1$, only one or no binary variables in the set can take the value one, ensuring that the m^{th} predictor does not repeat in more than one index. In other words, if the j^{th} element of \mathbf{Z}_m , z_{jm} , takes the value one, none of the other elements in the set \mathbf{Z}_m can take the value one, indicating that the m^{th} predictor enters into the j^{th} index in the model. On the other hand, if all the elements of the set \mathbf{Z}_m are zero, then the m^{th} predictor will be dropped out from the model.

Thus, our main contribution in this paper is two-fold. Firstly, we propose a novel algorithm to objectively estimate a semi-parametric additive index model, while contributing towards an estimated model with a higher forecasting accuracy. Secondly, the proposed variable selection methodology will contribute towards estimating a parsimonious model in a high-dimensional setting, even if the required domain knowledge for selecting the best set of predictors is unavailable.

3.4 Estimation Algorithm

In this section, we show how to efficiently find a minimiser for the problem given in Equation 8. Since the number of indices p , the vector of index coefficients α , as well as the set of nonparametric functions g are unknown, it is mathematically impossible to solve the above MIQP given in Equation 8 directly. Hence, we propose an iterative algorithm to solve the problem.

3.4.1 Initialising the Index Structure and Index Coefficients

Since the number of indices, the split of the predictors among indices as well as the index coefficients are not pre-specified, first, we need to provide an initialisation for the index structure (i.e. number of indices (p) and the split of predictors among indices) and the index coefficients (α) of the model for start solving the MIQP given in Equation 8.

Based on our experiments on the new algorithm, we propose five alternative methods for initialising the SMI Model as follows.

1. “PPR” - Projection Pursuit Regression Based Initialisation:

As discussed in Section 2.2, Projection Pursuit Regression model is a multiple index model, where each index consists of all the available predictors. Since in SMI Model we assume that there are no overlapping indices, it is impossible to use an estimated PPR model directly as a starting model for the algorithm. Thus, we follow the steps presented below to come up with a feasible initialisation for the index structure and the coefficients.

- i. Scale all the variables of the data set by dividing each variable by its standard deviation (so that it is possible to compare the estimated coefficients among predictors).
- ii. Fit a PPR model and obtain estimated index coefficients. (The user can decide the number of indices to be estimated (num_ind); we use $num_ind = 5$ as the default value.)
- iii. Calculate a threshold as

$$threshold = \max(PPR \text{ coefficients}) * 0.1$$

- iv. Assign zero for all the coefficients that are lower than the calculated threshold.
- v. If any predictor appears in more than one index, assign that predictor to the index in which that particular predictor has the maximum coefficient (among the coefficients corresponding to that predictor in all the indices), and make the coefficients corresponding to that predictor to be zero in all the other indices.
- vi. After performing the above steps i-v, if any of the originally estimated indices have all zero coefficients, such an index will be dropped out of the model.

Now, the index structure and the index coefficients obtained through the above steps are considered to a feasible initialisation for the SMI Model algorithm.

Furthermore, it is important to state here that once the optimal SMI Model is obtained through the proposed algorithm, each index coefficient will be back-transformed into the original scale of the corresponding predictor variable to roll back the effect of scaling at the beginning.

2. “Additive” - Nonparametric Additive Model Based Initialisation:

As mentioned previously in Section 3.1, the nonparametric additive model is a special case of SMI Model, where the number of indices equals the number of predictors entering indices ($p = q$) (i.e. number of predictors in each index = 1). Hence, it is a feasible starting point for the SMI Model algorithm. (Notice that in this case, we consider the index coefficient corresponding to each predictor to be 1.)

3. “Linear” - Linear Regression Based Initialisation:

In this option, we first regress the response variable on the predictors using a multiple linear regression. Then, we construct a single index (i.e. $p = 1$) using the estimated regression coefficients as the index coefficients of the predictors. Since Single Index Model is also a special case of the SMI Model, this will be a feasible starting point.

4. “Multiple” - Selecting an Initial Model by Comparing Multiple Models:

Through our experiments on the new algorithm, we identified that in some situations, the final optimised SMI Model often changes based on the initialisation that we provide to the algorithm. Hence, through this initialisation option, we consider a number of different models as initial models, optimise the SMI Model for each of those initial models, and pick the initial model that results in the SMI Model with the lowest loss for the MIP problem.

Here, the user can decide on the number of models to be considered (*num_models*) as well as the the number of indices to be considered in all the models (*num_ind* - same for all *num_models*). We use *num_models* = 5 and *num_ind* = 5 as default values.

5. “User Input” - User Specified Initialisation:

As mentioned earlier, since the number of indices, the split of the predictors among indices as well as the index coefficients are unknown, and will be optimally estimated by the proposed algorithm, theoretically, it is possible to start the algorithm at any given random initialisation of the model. Hence, this final option allow the user to provide their desired initialisation for the algorithm by specifying the number of indices, the split of predictors among the indices and the initial index coefficients.

In other words, this option provides the freedom for the user to utilise their domain expertise or prior knowledge in initialising the algorithm.

In all of the above initialisation options, once the estimate for α is obtained, the estimated initial index coefficient vector of each index $\hat{\alpha}_j = \alpha_{j,init}$ is scaled to have unit norm to ensure identifiability.

The characteristics and the performance of the proposed algorithm based on each of these initialisation options, vary depending on the nature of the application. Some initial experimental insights regarding the same will be discussed in more detail in Section 4.

3.4.2 Estimating Nonlinear Functions

Once we have an estimate for α , estimating the SMI Model is equivalent to estimating a GAM as

$$y_i = \beta_0 + \sum_{j=1}^p g_j(\hat{h}_{ij}) + \sum_{k=1}^d f_k(w_{ik}) + \theta^T \mathbf{u}_i + \varepsilon_i, \quad i = 1, \dots, n,$$

where y_i is taken as the response, and the estimated indices $\hat{h}_{ij} = \hat{\alpha}_j^T \mathbf{x}_i$, and the additional covariates that are not entering any index are taken as predictors.

The R packages *mgcv* (Wood 2011), and *gam* (Hastie 2023), for example, can be used to fit GAMs.

3.4.3 Updating the Index Structure and Index Coefficients

3.4.3.1 Updating Index Coefficients We estimate the new value of index coefficients α_{new} through a MIQP:

$$\begin{aligned} \min_{\alpha_{new}, z_{new}} \quad & (\alpha_{new} - \alpha_{old})^T \mathbf{V}^T \mathbf{V} (\alpha_{new} - \alpha_{old}) - 2(\alpha_{new} - \alpha_{old})^T \mathbf{V}^T \mathbf{r} \\ & + \lambda_0 \sum_{j=1}^p \sum_{m=1}^q z_{jm(new)} + \lambda_2 \sum_{j=1}^p \sum_{m=1}^q \alpha_{jm(new)}^2 \\ \text{s.t.} \quad & |\alpha_{jm(new)}| \leq M z_{jm(new)} \quad \forall j, \forall m, \\ & z_{jm(new)} \in \{0, 1\}, \\ & \sum_{j=1}^p z_{jm(new)} \leq 1 \quad \forall m, \\ & j = 1, \dots, p, \quad m = 1, \dots, q, \end{aligned} \tag{9}$$

where α_{old} is the current value of α , and $z_{jm(new)}$ are the new set of binary variables to be estimated. \mathbf{V} is the matrix of partial derivatives of the right hand side of Equation 6, with respect

to α_j . The i^{th} line of V contains $[v_{i1}, \dots, v_{ip}]$, where $v_{ij} = x_i g_j'(h_{ij})$. The current residual vector, which contains $r_i = y_i - \beta_0 - \sum_{j=1}^p g_j(\alpha_{j(old)}^T x_i)$, is denoted by r .

Similar to the explanation given by Masselot et al. (2022), the MIQP objective function in above Equation 9 ignores the *Hessian* (or the matrix of second derivatives of Equation 6, with respect to α_j), and considers only the matrix of first derivatives, which is a *quasi-Newton* step. The *Quasi-Newton Method* is an alternative to the *Newton's Method* that avoids the calculation of the Hessian to circumvent its computational burden (Peng 2022). Therefore, the α updating step given in above Equation 9 is assured to be in a *descent direction*.

It is important to note that the additional covariates w_{ik} and u_i do not step in to the process of updating α_j , because they are constants with respect to α_j , and thus they disappear from V .

Moreover, when α_{new} is obtained, if any of the estimated individual index coefficient vectors $\alpha_{j(new)}$ contains all zeros (i.e. zero index), such indices will be dropped out from the model.

Furthermore, similar to Section 3.4.1, once the new estimate α_{new} is obtained, we scale each estimated index coefficient vector $\hat{\alpha}_j = \alpha_{j(new)}$ to have unit norm.

3.4.3.2 Estimating Nonlinear Functions Similar to Section 3.4.2, once we obtain the new estimate α_{new} for α , we estimate a GAM taking y_i as the response, and the new estimated indices $\hat{h}_{ij(new)} = \hat{\alpha}_{j(new)}^T x_i$, and the additional covariates that are not entering any index as predictors.

3.4.3.3 Iterating We iterate the above steps described in Section 3.4.3.1 and Section 3.4.3.2 until the reduction ratio of the objective (loss) function value of the MIQP given in Equation 8, calculated between two successive iterations, reaches a pre-specified convergence tolerance, or the loss increases for three consecutive iterations. Alternatively, it is also possible to stop the iterations when a pre-specified maximum number of iterations is reached.

Here, to obtain an estimated model with the best possible forecasting accuracy, it is important to select appropriate values for the non-negative penalty parameters λ_0 and λ_2 . One possible way to do this is to estimate the model over a grid of possible values for λ_0 and λ_2 , and then select the combination that yields the lowest loss function value.

Moreover, it is also crucial to choose a suitable value for the big-M parameter, as the strength of the MIP formulation depends on the choice of a good lower bound (Bertsimas, King & Mazumder 2016). According to Hazimeh, Mazumder & Radchenko (2023), several methods have been used to select M in practice. For a description on estimating M in a linear regression setting, refer to Bertsimas, King & Mazumder (2016).

Additionally, the choice of convergence tolerance and the maximum number of iterations will depend on the nature of the problem/data to which the algorithm is applied. In the empirical applications presented in Section 5, we have used a convergence tolerance of 0.001, and 50 maximum iterations, where the iterations are terminated on whichever is reached first.

3.4.3.4 Updating Index Structure Once the initially chosen model is optimised through above step in Section 3.4.3.3, next, we consider changing the index structure of the model to exploit any benefits in terms of further minimising the loss function of the MIQP given in Equation 8.

Since reducing number of indices of the model automatically happens through dropping indices corresponding to any zero indices at each optimisation iteration, in this step, we consider adding indices into the current model.

Thus, once we obtain the optimal model at the step in Section 3.4.3.3 (say the number of indices in the model = p), we check for any predictors that are dropped out from the model. If any such predictors are available, we construct a new index comprising of those dropped out predictors and add that to the current model. Then the step in Section 3.4.3.3 is performed again for the new model with, say $p + 1$ indices.

If no predictors were dropped out from the first model (obtained from Section 3.4.3.3), then the algorithm terminates at this point, and the optimal model obtained from Section 3.4.3.3 will be the final SMI Model.

3.4.3.5 Iterating We iteratively perform the step in Section 3.4.3.4, increasing the number of indices of the existing model by one in each iteration. At the end of each iteration we check for the following termination criteria:

- *If the number of indices in the current model **equals** the number of indices in the new model:*
 - *If the absolute difference between the index coefficients in the current model and the new model is **equal** up to a pre-specified tolerance:*
 - * *If the loss of the new model is **greater** than the current model:* Terminate the algorithm, and the current model is the final SMI Model;
 - Else
 - * *If the loss of the new model is **lower** than the current model:* Terminate the algorithm, and the new model is the final SMI Model;
 - Else

- If the absolute difference between the index coefficients in the current model and the new model is **not equal** up to a pre-specified tolerance:
 - * If the loss of the new model is **greater** than the current model: Terminate the algorithm, and the current model is the final SMI Model;

Else

- If the number of indices in the current model **does not equal** the number of indices in the new model:
 - If the loss of the new model is **greater** than the current model: Terminate the algorithm, and the current model is the final SMI Model;

Else

- If the number of indices of the current model equals the number of predictors (entering indices) (i.e. maximum possible number of indices id reached): Terminate the algorithm, and the current model is the final SMI Model.

The following *Algorithm 1* summarises the key steps of the SMI Model algorithm.

Algorithm 1: SMI Model Algorithm

1. Initialise index structure and α :
Initialise p , split of the predictors among indices, and obtain α_{init} using one of the five methods in Section 3.4.1, and scale each $\hat{\alpha}_j = \alpha_{j,init}$ to have unit norm
2. Estimate g_j s:
Estimate g_j s using a GAM taking y_i as the response, $\hat{h}_{ij} = \hat{\alpha}_j^T x_i$ s as predictors
3. Update index structure and α :
 - i. Estimate the new value α_{new} through the MIQP in Equation 9, and scale each $\hat{\alpha}_j = \alpha_{j(new)}$ to have unit norm.
 - ii. Estimate new g_j s as in step 2.
 - iii. Iterate above steps i and ii until a convergence tolerance is reached OR loss increases for three consecutive iterations OR maximum number of iterations is reached.
 - iv. Update index structure:

- a. If any predictors are dropped out from the model in above step iii, add a new index consisting of the dropped out predictors, and optimise the model through step iii.
 - b. If no predictors are dropped out from the model in above step iii, terminate the algorithm.
- v. Iterate above step iv, while increasing p by 1 in each iteration, until one of the following terminations criteria is reached.
- **If** (No.of indices in current model = No.of indices in new model):
 - **If** (abs(current index coefficients – new index coefficients) \leq tolerance):
 - * **If** (loss of new model \geq loss of current model): **Terminate** - current model is the final SMI Model;
 - Else**
 - * **If** (loss of new model $<$ loss of current model): **Terminate** - new model is the final SMI Model;
 - Else**
 - **If** (abs(current index coefficients – new index coefficients) $>$ tolerance):
 - * **If** (loss of new model \geq loss of current model): **Terminate** - current model is the final SMI Model;
 - Else**
 - **If** (No.of indices in current model \neq No.of indices in new model):
 - **If** (loss of new model \geq loss of current model): **Terminate** - current model is the final SMI Model;
 - Else**
 - **If** (No. of indices in current model = q): **Terminate** - current model is the final SMI Model.

4 Simulation Experiment

This section presents the results of a simple simulation experiment conducted to demonstrate the performance and characteristics of the proposed SMI Model algorithm. Particularly, we try to investigate how the estimated SMI Model varies depending on the initialisation (as discussed in Section 3.4.1) used.

4.1 Data Generation

Generating predictor variables:

First, we generate two series each of length 1205: x , from random uniform distribution with the support $[0, 1]$, and z , from random normal distribution with mean = 5, and standard deviation = 2. Next, we construct lagged series up to 5th lag of both x and z . These current and lagged series of x and z (i.e. $x = x_0, x_1, \dots, x_5, z = z_0, z_1, \dots, z_5$) were taken as predictors in the simulation experiment.

Generating response variables:

We generated two response variables y_1 and y_2 , with two different index structures: “single-index” and “2-index”, and added a random normal noise component of two different strengths as follows:

- Low noise level - $N(\mu = 0, \sigma = 0.1)$:

$$y_1 = (0.9 * x_0 + 0.6 * x_1 + 0.45 * x_3)^3 + N(\mu = 0, \sigma = 0.1)$$

$$y_2 = (0.9 * x_0 + 0.6 * x_1 + 0.45 * x_3)^3 + (0.35 * x_2 + 0.7 * x_5)^2 + N(\mu = 0, \sigma = 0.1)$$

- High noise level - $N(\mu = 0, \sigma = 0.5)$:

$$y_1 = (0.9 * x_0 + 0.6 * x_1 + 0.45 * x_3)^3 + N(\mu = 0, \sigma = 0.5)$$

$$y_2 = (0.9 * x_0 + 0.6 * x_1 + 0.45 * x_3)^3 + (0.35 * x_2 + 0.7 * x_5)^2 + N(\mu = 0, \sigma = 0.5)$$

Hence, notice that the response y_1 is constructed using a single index consisting of the predictor variables x_0, x_1 , and x_3 , whereas the other response y_2 is constructed using two indices, where the first index consists of the predictors x_0, x_1 , and x_3 , and the second index consists of x_2 and x_5 . The variable x_4 or none of the z variables were used in generating y_1 and y_2 .

Once the data set is constructed, the first five observations are discarded due to the missing values introduced by the lagged variables, which leaves a data set of 1200 observations. We use the first 1000 observations as the training set, whereas the remaining 200 observations are kept aside as the test set for evaluating the forecast error of the estimated models.

4.2 Experiment setup

In this section, we explain the set up of the simulation experiment conducted.

We estimated SMI Models through the proposed algorithm for each of the two response variables (the two “true models”), using three different sets of predictors as inputs. Our aim was to

examine whether the algorithm correctly picks the relevant predictor variables (and drops the irrelevant predictors), and estimate the correct index structure of the true model.

The three different sets of predictors considered are as follows:

1. All “ x ” variables (i.e. only the “relevant” predictors as the responses, y_1 and y_2 , were constructed using only “ x ” variables);
2. All “ x ” variables and all “ z ” variables (i.e. “relevant” predictors plus “irrelevant” predictors); and
3. A part of “ x ” variables (x_0, x_1 and x_2) and all “ z ” variables (i.e. some “relevant” predictors plus “irrelevant” predictors).

We applied the proposed SMI Model algorithm using each of the above predictor combinations, for each of the two variations of the responses in terms of the noise level. Moreover, we also considered each of the first four initialisation options that we discussed in Section 3.4.1, for each of the two responses. For “PPR” and “Multiple” options, we used the default values of the required parameters as stated in Section 3.4.1. (Moreover, we used $M = 10$, a convergence tolerance of 0.001 and 50 maximum iterations in step 3.3 (Section 3.4.3.3), and a convergence tolerance for coefficients of 0.001 in step 3.5 (Section 3.4.3.5) in estimating all SMI Models.)

4.3 Results

We summarise the results of the simulation experiment in Table 1. Here, under the four columns corresponding to each of the four initialisation options considered, we indicate the index structure (i.e. the number of indices and the split of the variables among indices) estimated by the proposed algorithm, for each of the different combinations investigated based on response, input predictors, and noise levels.

Firstly, when the low noise level is considered, for the first true model y_1 , when either “all x ” or “all x + all z ” predictors are provided, all four initialisation options estimate the correct index structure of the model, while selecting the correct set of predictors and dropping out the irrelevant predictors. When only a part of the “ x ” variables are provided, the models estimated using all four initialisations have selected some noise variables. This indicates that when the available predictors are not sufficient to capture the signal of the data, the algorithm might select irrelevant variables to make up for the missing signal.

Table 1: *Simulation experiment results*

True Model	Predictors	PPR	Additive	Linear	Multiple
Low noise level					
y_1	all x	1 index	1 index	1 index	1 index
		(x_0, x_1, x_3)	(x_0, x_1, x_3)	(x_0, x_1, x_3)	(x_0, x_1, x_3)
	all x + all z	1 index	1 index	1 index	1 index
		(x_0, x_1, x_3)	(x_0, x_1, x_3)	(x_0, x_1, x_3)	(x_0, x_1, x_3)
	some x + all z	1 index	3 indices	1 index	1 index
		(x_0, x_1, z_2, z_4)	$(x_0, x_1) (z_4) (z_1)$	(x_0, x_1, z_2, z_4)	(x_0, x_1, z_2, z_4)
y_2	all x	2 indices	2 indices	1 index	2 indices
		$(x_0, x_1, x_3) (x_2, x_5)$	$(x_0, x_1, x_3) (x_2, x_5)$	$(x_0, x_1, x_2, x_3, x_5)$	$(x_0, x_1, x_3) (x_2, x_5)$
	all x + all z	2 indices	2 indices	1 index	2 indices
		$(x_0, x_1, x_3) (x_2, x_5)$	$(x_0, x_1, x_3) (x_2, x_5)$	$(x_0, x_1, x_2, x_3, x_5)$	$(x_0, x_1, x_3) (x_2, x_5)$
	some x + all z	3 indices	2 indices	1 index	2 indices
		$(x_0, x_1, z_2) (x_2) (z_3, z_4)$	$(x_0, x_1, z_4) (x_2)$	(x_0, x_1, x_2, z_2)	$(x_0, x_1) (x_2, z_2, z_3)$
High noise level					
y_1	all x	1 index	2 indices	1 index	1 index
		(x_0, x_1, x_3)	$(x_0, x_1) (x_3)$	(x_0, x_1, x_3)	(x_0, x_1, x_3)
	all x + all z	1 index	2 indices	1 index	2 indices
		(x_0, x_1, x_3)	$(x_0, x_1, x_3) (z_0)$	(x_0, x_1, x_3)	$(x_0, x_1, x_3) (z_0)$
	some x + all z	3 indices	3 indices	1 index	3 indices
		$(x_0, x_1) (z_1) (z_4)$	$(x_0, x_1) (z_1) (z_4)$	(x_0, x_1, z_2, z_4)	$(x_0, x_1) (z_0, z_4) (z_1)$
y_2	all x	3 indices	3 indices	2 indices	3 indices
		$(x_0, x_1, x_3) (x_2, x_5) (x_4)$	$(x_0, x_1, x_3) (x_2, x_5) (x_4)$	$(x_0, x_1, x_2, x_3, x_5) (x_4)$	$(x_0, x_1, x_3) (x_2, x_5) (x_4)$
	all x + all z	2 indices	2 indices	1 index	2 indices
		$(x_0, x_1, x_3) (x_2, x_5, z_1)$	$(x_0, x_1, x_3) (x_2, x_5, z_1)$	$(x_0, x_1, x_2, x_3, x_5, z_0)$	$(x_0, x_1, x_3) (x_2, x_5)$
	some x + all z	3 indices	2 indices	1 index	2 indices
		$(x_0, x_1, z_0, z_3) (x_2) (z_1, z_4, z_5)$	$(x_0, x_1, z_0, z_1, z_3, z_4) (x_2)$	$(x_0, x_1, x_2, z_0, z_3, z_4)$	$(x_0, x_1, z_0, z_1, z_3, z_4) (x_2)$

When y_2 with low noise level is considered, in both “all x” and “all x + all z” cases, except for the “Linear” initialisation, the algorithm has estimated the correct index structure with all the other three initialisation options. The “Linear” option has selected the correct variables into the model, however has not identified the 2-index structure. This suggests that it might be better to initialise the algorithm with a higher number of indices rather than with a lower number of indices. When only a part of “x” variables are provided, similar to the case of y_1 , models estimated for y_2 with all four initialisations, have selected noise variables.

Furthermore, we evaluated the forecasting error of the fitted models on the test set using Mean Squared Error (MSE). For y_1 , in both the cases “all x” and “all x + all z”, the models estimated using all four initialisations resulted in a test MSE of ≈ 0.01 , which is the random squared error of the true model. This confirms the accuracy with which the SMI Model algorithm estimated the index structure for y_1 . When only a part of “x” variables are provided, the test MSE increased to ≈ 0.23 , as the estimated models contain noise variables.

For y_2 , in both the cases “all x ” and “all x + all z ”, all the models estimated resulted in a test MSE of ≈ 0.16 . This is an interesting result as the test MSEs of the estimated models with correct index structure: “PPR”, “Additive”, and “Multiple”, are not very different to the test MSE of an estimated model with incorrect index structure, but with correct predictors: “Linear”. This suggests that the selection of the predictor variables is more important than determining the index structure of the model.

It is also worth mentioning here that in real-world forecasting problems, the true data generating process (DGP) is unknown, and we do not expect an estimated model to represent the true DGP. Therefore, as long as the estimated model exhibits good forecasting accuracy, the index structure of the estimated model is less important.

Similar to the case of y_1 , when only a part of “ x ” variables are provided, the test MSE of the estimated models for y_2 increased to ≈ 0.34 , probably due to the inclusion of noise variables.

Moreover, when the above test MSE values of the models estimated for y_2 are considered, in contrast to the case of y_1 , those values are higher than the random squared error of the true model. Intuitively, this might be due to the increased complexity of the model y_2 in comparison to y_1 , where the total estimation error of two nonlinear link functions (corresponding to the two indices) for y_2 , might be higher than the error of estimating a single nonlinear function for y_1 .

As expected, the accuracy with which the SMI Model algorithm estimates the index structure is in general lower with the high noise level, in comparison to the low noise level. In the case of y_1 , when “all x ” variables are provided, except for “Additive” option, all the other initialisations have correctly estimated the model. When “all x + all z ” variables are provided, only the “PPR” and “Linear” options have estimated the correct model, whereas “Additive” and “Multiple” options have included a noise variable. Similar to the low noise level case, when only a part of “ x ” variables are provided, all four initialisations have led the algorithm to estimate models with noise variables.

Next, when y_2 is considered at high noise level, when “all x ” variables are provided, the estimated models using all four initialisation options have included the variable x_4 (as an additional index), which should have been omitted. When “all x + all z ” variables are provided, the option “Multiple” has led the algorithm to the correct model, while all the other three options have included a noise variable. Similar to the previous cases, when only a part of “ x ” variables are provided, the models estimated using all four initialisation options have included multiple noise variables, probably to cover up for the missing signal variables.

In both the cases “all x ” and “all x + all z ”, all the models estimated for y_1 , except for “Additive” option in “all x ” case, resulted in a test MSE of ≈ 0.23 (which is slightly lower than the random squared error of the true model; this probably indicates a slight level of over-fitting), irrespective of the fact that in “all x + all z ” case, “Additive” and “Multiple” options included a noise variable. This is an indication of the effect of the low signal-to-noise ratio in the data. The model estimated using “Additive” option in “all x ” case has resulted in a slightly higher test MSE (≈ 0.27), probably due to the incorrect index structure. When only a part of “ x ” variables are provided, the test MSE has increased (≈ 0.47), which is expected.

Similar to previously discussed low noise level case, the estimated models for y_2 have resulted in higher test MSE values in comparison to the models for y_1 . When “all x ” and “all x + all z ” variables are provided, all the estimated models resulted in test MSE values in the range $\approx (0.35, 0.37)$, irrespective of the different index structure and predictor choices, which is again an indication of the low signal-to-noise ratio in the data. Due to missing relevant predictors (and the inclusion of noise predictors in the model), the estimated models in “some x + all z ” case resulted in test MSEs ≈ 0.57 .

In the above simulation experiment, we did not perform any tuning for the penalty parameters λ_0 and λ_2 , as through our experiments we observed that for this simple example, different values of penalty parameters have a negligible impact on the estimated models. We used the default values $\lambda_0 = 1$ and $\lambda_2 = 1$ in estimating all the models presented above in Table 1.

Finally, the above simulation study indicates that the choice of the initialisation depends on the data/application. Thus, the users of the algorithm are encouraged to follow a trial-and-error procedure for choosing the best initial model for a given application.

In the next section, we illustrate the forecasting performance of the SMI Model in two real-world applications.

5 Empirical Applications

(Major editing required.)

5.1 Forecasting Daily Mortality

We apply the SGAIM algorithm to a data set from Masselot et al. (2022), to forecast daily mortality based on heat exposure.

Studying the effects of various environmental exposures such as weather related variables, pollutants and man-made environmental conditions etc. on human health, is of significant importance in environmental epidemiology. Therefore, forecasting daily deaths taking heat related variables as predictors, and constructing interpretable indices of those predictors that reflect heat-related mortality risk, is an interesting application.

5.1.1 Description of the Data

For this analysis, we consider daily mortality and heat exposure data for the Metropolitan Area of Montreal, Province of Quebec, Canada, from 1990 to 2014, for the months June, July, and August (i.e. summer season). The daily all-cause mortality data were obtained from the National Institute of Public Health, Province of Quebec, while “DayMet” - a 1 km \times 1 km gridded data set (Thornton et al. 2021) was used to extract daily temperature and humidity data (Masselot et al. 2022).

Figure 1 shows the time plots of daily deaths during the summer for the years from 1990 to 1993. The series for only four years are plotted separately as a faceted grid for visual clarity.




Figure 1: *Daily mortality in summer in Montreal, Canada - from 1990 to 1993*

The three main predictors considered in this empirical study are maximum temperature, minimum temperature, and vapor pressure (to represent the level of humidity). The number of daily deaths are plotted against each of these predictors in Figure 2, Figure 3, and Figure 4 respectively,

where we can observe that the relationships between all these predictors and the response are slightly non-linear.




Figure 2: Daily mortality in summer (from 1990 to 2014) plotted against maximum temperature

5.1.2 Predictors Considered

1) Current maximum/minimum temperatures and lags:

In addition to current maximum and minimum temperatures, the temperature measurements up to 14 days prior (i.e. 0^{th} to 14^{th} lag) are also considered as predictors to the forecasting model, because it is obvious that not only the current temperature, but also the temperatures that were prevailing in the recent past can add up to the adverse level of heat exposure of a person.

2) Current vapor pressure and lags:

Similar to temperature variables, the current value and 14 lags of vapor pressure are considered as predictors, as a proxy to the level of humidity.

3) Calendar effects:

Finally, a couple of calendar variables; *day of the season (DOS)* and *Year*, are incorporated into the model to capture annual trend and seasonality, and also to control the autocorrelation in residuals, which is a common practice in environmental epidemiology (Masselot et al. 2022).




Figure 3: Daily mortality in summer (from 1990 to 2014) plotted against minimum temperature




Figure 4: Daily mortality in summer (from 1990 to 2014) plotted against vapor pressure

5.1.3 Modelling Framework

Maximum temperature lags, minimum temperature lags and vapor pressure lags are categorised into three groups of predictors, where we estimate an index for each of those groups using the proposed SGAIM algorithm.

The two calendar variables, *DOS* and *Year*, are included into the model as separate nonparametric components that do not enter any of the indices.

Hence, the relevant SGAIM can be expressed as

$$\begin{aligned} \text{Deaths} = & \beta_0 + g_1(\mathbf{Tmax_Lags} * \boldsymbol{\alpha}_1) + g_2(\mathbf{Tmin_Lags} * \boldsymbol{\alpha}_2) + g_3(\mathbf{Vp_Lags} * \boldsymbol{\alpha}_3) \\ & + f_1(\mathbf{DOS}) + f_2(\mathbf{Year}) + \boldsymbol{\varepsilon}, \end{aligned} \quad (10)$$

where

- **Deaths** is the vector containing the observations of number of daily deaths;
- **Tmax_Lags** is a matrix containing lags $0, \dots, 14$ of maximum temperature;
- **Tmin_Lags** is a matrix containing lags $0, \dots, 14$ of minimum temperature;
- **Vp_Lags** is a matrix containing lags $0, \dots, 14$ of vapor pressure;
- $\boldsymbol{\alpha}_j, j = 1, 2, 3$ are the index coefficient vectors of length 15 each;
- $g_j, j = 1, 2, 3, f_1$ and f_2 are unknown nonparametric functions;
- β_0 is the model intercept;
- $\boldsymbol{\varepsilon}$ is the vector of errors.

The data from 1990 to 2012 are used as the training set to estimate the model. The data corresponding to the three summer months of year 2013 are kept aside as the validation set, while the data of year 2014 are separated to be the test set to evaluate the forecasting performance of the estimated model.

Then we apply the proposed SGAIM algorithm to estimate the index coefficient vectors $\boldsymbol{\alpha}_j, j = 1, 2, 3$ (Equation 10). After obtaining the optimal values for $\boldsymbol{\alpha}_j$ s, we calculate the indices, and fit a GAM on the training data, taking the estimated indices and the calendar variables as predictors. Then the forecasting accuracy on the test set is evaluated using Mean Squared Error (MSE) and Mean Absolute Error (MAE).

5.1.4 Results

Point Forecasts:

We considered 14 lags of each heat-related predictor variable in the model. However, the CGAIM fitted in Masselot et al. (2022) has considered only up to the 2nd lag of the each predictor. Hence, for comparison purposes, we fitted two sets of models; one set of models with 14 lags and the other set with only 2 lags.

Similar to the previous application, we tuned the penalty parameters λ_0 and λ_2 , over ranges of integers from 1 to 15, and 0 to 15 respectively, on the validation set.

Case 1: SGAIM with 14 lags of each predictor

The combination ($\lambda_0 = 15, \lambda_2 = 0$) (i.e. without ridge penalty) yielded the lowest MSE and MAE on the validation set (*SGAIM (15, 0)*), where 4 maximum temperature lags (0, 1, 2, 3) were selected by the algorithm for the *maximum temperature index*, 2 lags of minimum temperature (2, 14) was selected for the *minimum temperature index*, and 3 lags of vapor pressure (1, 2, 12) were selected for the *vapor pressure index*. Hence, the algorithm selected only 9 variables for the estimated *SGAIM (15, 0)*, out of the total 45 predictors entering indices. Similar to the previous application, we used $M = 10$, a convergence tolerance of 0.001 and 50 maximum iterations in estimating all SGAIMs.

We evaluated the forecasting error of the model selected using two different subsets of the original test set:

1. *Test Set 1*: original test set of 3 months (June, July and August of 2014); and
2. *Test Set 2*: a test set of 1 month (June 2014).

Note that similar to the previous application of electricity demand forecasting, we assumed that the future values of the maximum/minimum temperatures and vapor pressure are known to use in the forecasting model.

The MSE and MAE values obtained on the two variations of the test set for *SGAIM (15, 0)* are presented in Table 2. The forecasting errors of a CGAIM with the number of predictor lags increased to 14, are also presented for comparison purposes. Here, following Masselot et al. (2022), the index coefficients of the CGAIM were constrained to be positive and decreasing, and the nonparametric link functions were constrained to be monotonically increasing. We also fitted an unconstrained GAIM (without variable selection), where the results are presented in

Table 2. The actual series of number of deaths and the predicted series from the $SGAIM(15, 0)$ and the benchmark models on *Test Set 2* are plotted in Figure 5 for further comparison.

Table 2: Daily mortality forecasting (with 14 lags of predictors) - Out-of-sample point forecast results

Model	Predictors	Test Set 1		Test Set 2	
		MSE	MAE	MSE	MAE
$SGAIM(15, 0)$	11	87.577	7.160	102.483	7.870
CGAIM	47	85.679	7.208	101.689	8.032
Unconstrained GAIM	47	81.740	6.920	103.036	8.066

According to Table 2, in terms of the test MSE, $SGAIM(15, 0)$ is unable to outperform either the constrained or unconstrained GAIMs in *Test Set 1*. $SGAIM(15, 0)$ reports a slightly lower test MSE than the unconstrained GAIM, and a slightly lower test MAE in comparison to both CGAIM and unconstrained GAIM in *Test Set 2*, but the reductions are not of a considerable magnitude. However, note that $SGAIM(15, 0)$ achieves a forecasting performance comparable to the benchmark models with a considerably lesser number of predictors.




Figure 5: Actual number of deaths vs. predicted number of deaths from “ $SGAIM(15, 0)$ ” and benchmark models for *Test Set 2*

Case 2: $SGAIM$ with 2 lags of each predictor

For the $SGAIM$ with only up to 2 lags of the predictors, the combination ($\lambda_0 = 1, \lambda_2 = 0$) (i.e. without ridge penalty) yielded the lowest MSE on the validation set ($SGAIM(1, 0)$). This

model selected all the 3 lags (0, 1, 2) of maximum temperature for the *maximum temperature index*, only the current value of minimum temperature for the *minimum temperature index*, and lags 1 and 2 of vapor pressure for the *vapor pressure index*. Hence, the algorithm selected 6 predictors out of the total 9 predictors entering indices.

We evaluated the forecasting error of the estimated $SGAIM(1, 0)$ using the same two different subsets of the original test set, where the results are reported in Table 3. The forecasting errors of CGAIM and unconstrained GAIM with only up to 2 lags for each predictor, are also presented as benchmark models. The actual series of number of deaths and the predicted series from the $SGAIM(1, 0)$ and the benchmark models on *Test Set 2* are plotted in Figure 6 for further comparison.

Table 3: Daily mortality forecasting (with 2 lags of predictors) - Out-of-sample point forecast results

Model	Predictors	Test Set 1		Test Set 2	
		MSE	MAE	MSE	MAE
SGAIM (1, 0)	8	85.892	7.186	103.321	7.999
CGAIM	11	83.625	7.011	101.689	8.032
Unconstrained GAIM	11	83.458	7.066	103.036	8.066

According to Table 3, the SGAIM fitted considering only 2 lags of each predictor too was not able to outperform the corresponding CGAIM or unconstrained GAIM in terms of MSE. Similar to above *Case 1*, the test MAE reported by the $SGAIM(1, 0)$ is slightly lower than the CGAIM and unconstrained GAIM in *Test Set 2*, but again the reduction is not considerably large.

Both the empirical applications presented above were performed using *R statistical software* (R Core Team 2023), and the *Rstudio* integrated development environment (IDE) (Posit team 2023). We used the commercial MIP solver “*Gurobi*” (Gurobi Optimization, LLC 2023) to solve the MIQPs related to the proposed SGAIM algorithm, through the *Gurobi plug-in (ROI.plugin.gurobi)* (Schwendinger 2023) available from the *R Optimization Infrastructure (ROI)* (Hornik et al. 2023; Theußl, Schwendinger & Hornik 2020) package. Furthermore, the GAMs were fitted using the R package *mgcv* (v1.8.42) (Wood 2011).

6 Conclusions

(To be completed.)

Acknowledgement




Figure 6: Actual number of deaths vs. predicted number of deaths from “SGAIM (1, 0)” and benchmark models for Test Set 2

We thank Professor Louise Ryan for joining the discussions during the initial stage of the project, and for her valuable comments and feedback on this research work.

References

- Bakker, M & F Schaars (2019). Solving groundwater flow problems with time series analysis: you may not even need another model. *Groundwater* **57**(6), 826–833.
- Bellman, R (1957). *Dynamic Programming*. Princeton, New Jersey: Princeton University Press.
- Bertsimas, D, A King & R Mazumder (2016). Best subset selection via a modern optimization lens. *Annals of Statistics* **44**(2), 813–852.
- Boggs, PT & JW Tolle (1995). Sequential Quadratic Programming. *Acta Numerica* **4**, 1–51.
- Boyd, SP & L Vandenberghe (2004). *Convex Optimization*. Cambridge University Press.
- Breiman, L (1995). Better Subset Regression Using the Nonnegative Garrote. *Technometrics* **37**(4), 373–384.
- Efron, B, T Hastie, I Johnstone & R Tibshirani (2004). Least Angle Regression. *Annals of Statistics* **32**(2), 407–499.
- Fan, J & R Li (2001). Variable Selection via Nonconcave Penalized Likelihood and its Oracle Properties. *Journal of the American Statistical Association* **96**(456), 1348–1360.

- Fan, S & RJ Hyndman (2012). Short-term load forecasting based on a semi-parametric additive model. *IEEE Transactions on Power Systems* **27**(1), 134–141.
- Frank, IE & JH Friedman (1993). A Statistical View of Some Chemometrics Regression Tools. *Technometrics* **35**(2), 109–135.
- Friedman, JH & JW Tukey (1974). A Projection Pursuit Algorithm for Exploratory Data Analysis. *IEEE Transactions on Computers* **C-23**(9), 881–890.
- Friedman, JH & W Stuetzle (1981). Projection Pursuit Regression. *Journal of American Statistical Association* **76**(376), 817–823.
- Guo, J, M Tang, M Tian & K Zhu (2013). Variable selection in high-dimensional partially linear additive models for composite quantile regression. *Computational Statistics and Data Analysis* **65**, 56–67.
- Gurobi Optimization, LLC (2023). *Gurobi Optimizer Reference Manual*. <https://www.gurobi.com>.
- Härdle, W, P Hall & H Ichimura (1993). Optimal Smoothing in Single-Index Models. *Annals of Statistics* **21**(1), 157–178.
- Hastie, T (2023). *gam: Generalized Additive Models*. R package version 1.22-2. <https://CRAN.R-project.org/package=gam>.
- Hazimeh, H & R Mazumder (2020). Fast Best Subset Selection: Coordinate Descent and Local Combinatorial Optimization Algorithms. *Operations Research* **68**(5), 1517–1537.
- Hazimeh, H, R Mazumder & P Radchenko (2023). Grouped variable selection with discrete optimization: Computational and statistical perspectives. *Annals of Statistics* **51**(1), 1–32.
- Ho, CC, LJ Chen & JS Hwang (2020). Estimating ground-level PM2.5 levels in Taiwan using data from air quality monitoring stations and high coverage of microsensors. *Environmental Pollution* **264**, 114810.
- Hornik, K, D Meyer, F Schwendinger & S Theussl (2023). *ROI: R Optimization Infrastructure*. R package version 1.0-1. <https://CRAN.R-project.org/package=ROI>.
- Huang, J, JL Horowitz & F Wei (2010). Variable Selection in Nonparametric Additive Models. *Annals of Statistics* **38**(4), 2282–2313.
- Huber, PJ (1985). Projection Pursuit. *Annals of Statistics* **13**(2), 435–475.
- Hyndman, RJ & S Fan (2010). Density forecasting for long-term peak electricity demand. *IEEE Transactions on Power Systems* **25**(2), 1142–1153.
- Ibrahim, S, R Mazumder, P Radchenko & E Ben-David (2022). “Predicting Census Survey Response Rates via Interpretable Nonparametric Additive Models with Structured Interactions”. <https://arxiv.org/abs/2108.11328>.

- Konzen, E & FA Ziegelmann (2016). LASSO-type penalties for covariate selection and forecasting in time series. *Journal of Forecasting* **35**(7), 592–612.
- Kruskal, JB (1969). “Toward a practical method which helps uncover the structure of a set of multivariate observations by finding the linear transformation which optimizes a new “index of condensation””. In: *Statistical Computation*. Ed. by Roy C. Milton and John A. Nelder. Academic Press, pp.427–440.
- Land, AH & AG Doig (1960). An Automatic Method of Solving Discrete Programming Problems. *Econometrica* **28**(3), 497–520.
- Lian, H (2012). Variable selection in high-dimensional partly linear additive models. *Journal of Nonparametric Statistics* **24**(4), 825–839.
- Little, JDC, KG Murty, DW Sweeney & C Karel (1963). An algorithm for the traveling salesman problem. *Operations Research* **11**(6), 972–989.
- Liu, X, L Wang & H Liang (2011). Estimation and Variable Selection for Semiparametric Additive Partial Linear Models (SS-09-140). *Statistica Sinica* **21**(3), 1225–1248.
- Masselot, P, F Chebana, C Campagna, É Lavigne, TBMJ Ouarda & P Gosselin (2022). Constrained groupwise additive index models. *Biostatistics* **00**(00), 1–19.
- Mazumder, R, P Radchenko & A Dedieuc (2022). Subset Selection with Shrinkage: Sparse Linear Modeling When the SNR Is Low. *Operations Research*, 1–19.
- Nocedal, J & SJ Wright (2006). *Numerical Optimization*. 2nd. Springer Series in Operations Research and Financial Engineering. Springer New York, NY.
- Park, H & F Sakaori (2013). Lag weighted lasso for time series model. *Computational Statistics* **28**(2), 493–504.
- Peng, RD (2022). *Advanced Statistical Computing*. <https://bookdown.org/rdpeng/advstatcomp/>. Accessed: 2023-5-19.
- Peterson, TJ & AW Western (2014). Nonlinear time-series modeling of unconfined groundwater head. *Water Resources Research* **50**(10), 8330–8355.
- Posit team (2023). *RStudio: Integrated Development Environment for R*. Posit Software, PBC. Boston, MA. <http://www.posit.co/>.
- R Core Team (2023). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing. Vienna, Austria. <https://www.R-project.org/>.
- Radchenko, P (2015). High dimensional single index models. *Journal of Multivariate Analysis* **139**, 266–282.
- Rajaei, T, H Ebrahimi & V Nourani (2019). A review of the artificial intelligence methods in groundwater level modeling. *Journal of Hydrology* **572**, 336–351.

- Ravindra, K, P Rattan, S Mor & AN Aggarwal (2019). Generalized additive models: Building evidence of air pollution, climate change and human health. *Environment International* **132**, 104987.
- Schwendinger, F (2023). *ROI.plugin.gurobi: 'Gurobi' Plug-in for the 'R' Optimization Infrastructure*. R package version 0.4-0. <http://r-forge.r-project.org/projects/roi>.
- Simon, N, J Friedman, T Hastie & R Tibshirani (2013). A Sparse-Group Lasso. *Journal of Computational and Graphical Statistics* **22**(2), 231–245.
- Stoker, TM (1986). Consistent Estimation of Scaled Coefficients. *Econometrica* **54**(6), 1461–1481.
- Stone, CJ (1982). Optimal global rates of convergence for nonparametric regression. *Annals of Statistics* **10**(4), 1040–1053.
- Theußl, S, F Schwendinger & K Hornik (2020). ROI: An Extensible R Optimization Infrastructure. *Journal of Statistical Software* **94**, 1–64.
- Thornton, PE, R Shrestha, M Thornton, SC Kao, Y Wei & BE Wilson (2021). Gridded daily weather data for North America with comprehensive uncertainty quantification. *Scientific Data* **8**(1), 190.
- Tibshirani, R (1996). Regression Shrinkage and Selection via the Lasso. *Journal of the Royal Statistical Society (Series B)* **58**(1), 267–288.
- Tibshirani, R & X Suo (2016). An Ordered Lasso and Sparse Time-Lagged Regression. *Technometrics* **58**(4), 415–423.
- Wang, H, L Guodong & CL Tsai (2007). Regression Coefficient and Autoregressive Order Shrinkage and Selection Via the Lasso. *Journal of the Royal Statistical Society (Series B)* **69**(1), 63–78.
- Wang, L, L Xue, A Qu & H Liang (2014). Estimation and model selection in generalized additive partial linear models for correlated data with diverging number of covariates. *Annals of Statistics* **42**(2), 592–624.
- Wang, T, P Xu & L Zhu (2015). Variable selection and estimation for semi-parametric multiple-index models. *Bernoulli* **21**(1), 242–275.
- Wang, T, J Zhang, H Liang & L Zhu (2015). Estimation of a Groupwise Additive Multiple-Index Model and its Applications. *Statistica Sinica* **25**, 551–566.
- Wood, SN (2017). *Generalized Additive Models: An Introduction with R*. 2nd. Chapman & Hall/CRC.
- Wood, SN (2011). Fast stable restricted maximum likelihood and marginal likelihood estimation of semiparametric generalized linear models. *Journal of the Royal Statistical Society (Series B)* **73**(1), 3–36.

- Yuan, M & Y Lin (2006). Model selection and estimation in regression with grouped variables. *Journal of the Royal Statistical Society (Series B)* **68**(1), 49–67.
- Zhang, CH (2010). Nearly Unbiased Variable Selection under Minimax Concave Penalty. *Annals of Statistics* **38**(2), 894–942.
- Zhang, X, L Liang, X Tang & HY Shum (2008). L1 regularized projection pursuit for additive model learning. In: *2008 IEEE Conference on Computer Vision and Pattern Recognition*, pp.1–8.
- Zou, H (2006). The Adaptive Lasso and its Oracle Properties. *Journal of the American Statistical Association* **101**(476), 1418–1429.